

ML0101EN-Reg-Polynomial-Regression-Co2-py-v1

May 10, 2019

Polynomial Regression

About this Notebook

In this notebook, we learn how to use scikit-learn for Polynomial regression. We download a dataset that is related to fuel consumption and Carbon dioxide emission of cars. Then, we split our data into training and test sets, create a model using training set, evaluate our model using test set, and finally use model to predict unknown value.

Table of contents

```
<ol>
  <li><a href="#download_data">Downloading Data</a></li>
  <li><a href="#polynomial_regression">Polynomial regression</a></li>
  <li><a href="#evaluation">Evaluation</a></li>
  <li><a href="#practice">Practice</a></li>
</ol>
```

0.0.1 Importing Needed packages

```
In [1]: import matplotlib.pyplot as plt
import pandas as pd
import pylab as pl
import numpy as np
%matplotlib inline
```

Downloading Data

To download the data, we will use `!wget` to download it from IBM Object Storage.

```
In [2]: !wget -O FuelConsumption.csv https://s3-api.us-gio.objectstorage.softlayer.net/cf-courses-data/CognitiveC
--2019-05-10 08:22:40-- https://s3-api.us-gio.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/ML0101E
Resolving s3-api.us-gio.objectstorage.softlayer.net (s3-api.us-gio.objectstorage.softlayer.net)... 67.228.254.193
Connecting to s3-api.us-gio.objectstorage.softlayer.net (s3-api.us-gio.objectstorage.softlayer.net)|67.228.254.193|:4
HTTP request sent, awaiting response... 200 OK
Length: 72629 (71K) [text/csv]
Saving to: FuelConsumption.csv
```

```
FuelConsumption.csv 100%[=====>] 70.93K --.-KB/s in 0.04s
```

```
2019-05-10 08:22:40 (1.63 MB/s) - FuelConsumption.csv saved [72629/72629]
```

Did you know? When it comes to Machine Learning, you will likely be working with large datasets. As a business, where can you host your data? IBM is offering a unique opportunity for businesses, with 10 Tb of IBM Cloud Object Storage: [Sign up now for free](#)

0.1 Understanding the Data

0.1.1 FuelConsumption.csv:

We have downloaded a fuel consumption dataset, FuelConsumption.csv, which contains model-specific fuel consumption ratings and estimated carbon dioxide emissions for new light-duty vehicles for retail sale in Canada. [Dataset source](#)

- **MODELYEAR** e.g. 2014
- **MAKE** e.g. Acura
- **MODEL** e.g. ILX
- **VEHICLE CLASS** e.g. SUV
- **ENGINE SIZE** e.g. 4.7
- **CYLINDERS** e.g 6
- **TRANSMISSION** e.g. A6
- **FUEL CONSUMPTION in CITY (L/100 km)** e.g. 9.9
- **FUEL CONSUMPTION in HWY (L/100 km)** e.g. 8.9
- **FUEL CONSUMPTION COMB (L/100 km)** e.g. 9.2
- **CO2 EMISSIONS (g/km)** e.g. 182 --> low --> 0

0.2 Reading the data in

```
In [3]: df = pd.read_csv("FuelConsumption.csv")
```

```
# take a look at the dataset
df.head()
```

```
Out[3]:  MODELYEAR  MAKE      MODEL VEHICLECLASS  ENGINESIZE  CYLINDERS \
0      2014  ACURA      ILX      COMPACT        2.0          4
1      2014  ACURA      ILX      COMPACT        2.4          4
2      2014  ACURA  ILX HYBRID      COMPACT        1.5          4
3      2014  ACURA    MDX 4WD  SUV - SMALL        3.5          6
4      2014  ACURA    RDX AWD  SUV - SMALL        3.5          6

      TRANSMISSION FUELTYPE  FUELCONSUMPTION_CITY  FUELCONSUMPTION_HWY \
0          AS5        Z          9.9          6.7
1          M6        Z         11.2          7.7
2          AV7        Z          6.0          5.8
3          AS6        Z         12.7          9.1
4          AS6        Z         12.1          8.7

      FUELCONSUMPTION_COMB  FUELCONSUMPTION_COMB_MPG  CO2EMISSIONS
```

0	8.5	33	196
1	9.6	29	221
2	5.9	48	136
3	11.1	25	255
4	10.6	27	244

Lets select some features that we want to use for regression.

```
In [4]: cdf = df[['ENGINE SIZE','CYLINDERS','FUELCONSUMPTION_COMB','CO2EMISSIONS']]
cdf.head(9)
```

```
Out[4]:  ENGINE SIZE  CYLINDERS  FUELCONSUMPTION_COMB  CO2EMISSIONS
0         2.0         4         8.5             196
1         2.4         4         9.6             221
2         1.5         4         5.9             136
3         3.5         6        11.1             255
4         3.5         6        10.6             244
5         3.5         6        10.0             230
6         3.5         6        10.1             232
7         3.7         6        11.1             255
8         3.7         6        11.6             267
```

Lets plot Emission values with respect to Engine size:

```
In [5]: plt.scatter(cdf.ENGINE SIZE, cdf.CO2EMISSIONS, color='blue')
plt.xlabel("Engine size")
plt.ylabel("Emission")
plt.show()
```

